

Chapter 8

IPC - II

Chapter Objectives

To understand concepts of IPC used over the duration of the course.

Objectives

For this chapter, the following are the objectives:

- Message Queues,
- Shared Memory,
- Semaphores

Slide #8-1

Notes

In this chapter, we examine System V (“System Five”) inter-process communication, mechanisms that are widely used.

The objective of this chapter is to provide an understanding of concepts on System V IPC that are used over the duration of the course.

Chapter Organization

1. **Objective:** Introduction to
 - message queues,
 - shared memory,
 - semaphores
2. **Description:** These are the various tools needed for systems programming
This chapter provides an introduction to the concepts used in course.
3. **Concepts Covered in Chapter:**
 - Introduction to various tools.
 - LINUX Notes
4. **Prior Knowledge:**
same as Chapter #1
5. **Teaching & Learning Strategy:**
Discussion questions are,
 - What are these tools for?
6. **Teaching Format:**
Theory + Homework Assignments
7. **Study Time:** 120 Minutes (Lecture & Theory)
+ ~45 minutes (Homework Assignments)
8. **Assessment:** Group Homework Assignments
9. **Homework Eval:** Group
10. **Chapter References:**

getopt

getopt

- Parse command line options
- Widely Used .. Portable

Slide #8-2

Notes**getopt – Basic Concepts**

UNIX provides a command line parser.

Code example (sourced from <http://www.die.net/doc/linux/man/man3/getopt.3.html>)

```
#include <stdio.h>      /* for printf */
#include <stdlib.h>     /* for exit */
#include <getopt.h>

int
main (int argc, char **argv) {
    int c;
    int digit_optind = 0;

    while (1) {
        int this_option_optind = optind ? optind : 1;
        int option_index = 0;
        static struct option long_options[] = {
            {"add", 1, 0, 0},
            {"append", 0, 0, 0},
        }
```

```
        {"delete", 1, 0, 0},
        {"verbose", 0, 0, 0},
        {"create", 1, 0, 'c'},
        {"file", 1, 0, 0},
        {0, 0, 0, 0}
    };

    c = getopt_long (argc, argv, "abc:d:012",
                    long_options, &option_index);
    if (c == -1)
        break;

    switch (c) {
    case 0:
        printf ("option %s",
                long_options[option_index].name);
        if (optarg)
            printf (" with arg %s", optarg);
        printf ("\n");
        break;

    case '0':
    case '1':
    case '2':
        if (digit_optind != 0 && digit_optind !=
            this_option_optind)
            printf ("digits occur in two different argv-
elements.\n");
        digit_optind = this_option_optind;
        printf ("option %c\n", c);
        break;

    case 'a':
        printf ("option a\n");
        break;

    case 'b':
        printf ("option b\n");
        break;

    case 'c':
        printf ("option c with value '%s'\n", optarg);
        break;

    case 'd':
        printf ("option d with value '%s'\n", optarg);
        break;

    case '?':
```

```
        break;

    default:
        printf ("?? getopt returned character code
0%o ??\n", c);
    }
}

if (optind < argc) {
    printf ("non-option ARGV-elements: ");
    while (optind < argc)
        printf ("%s ", argv[optind++]);
    printf ("\n");
}

exit (0);
}
```

System V -- IPC

System V - IPC

- Widely Used .. Portable
- IPC between unrelated processes

Slide #8-3

Notes**System V IPC – Basic Concepts**

1. UNIX provides a rich set of features for unrelated processes to communicate with each other. Inter-Process Communication (IPC), can be used to share data between unrelated processes running simultaneously.
 - 1.1. It is widely used, and is widely available.
 - 1.2. Portable across most UNIX systems.
2. Linux has System V IPC mechanisms for several years.

System V - IPC

- Shared Memory
- Message Queues
- Semaphores

Slide #8-4

System V IPC – Basic Concepts

1. The System V (SysV) UNIX specification describes three mechanisms for IPC, collectively referred to as SysV IPC:
 - 1.1. Message queues
 - 1.2. Semaphores
 - 1.3. Shared memory

System V - IPC

The Key

- Can use a hard coded value, or,
- Generate a key using `ftok()`

Slide #8-5

1. Hard Coded value, or,
2. Creating the key

```
key = ftok("./ch7.c", 'R');
```

Shared Memory

System V - Shared Memory

Memory

- Shared by multiple processes
- Persists till removed

Slide #8-6

Notes

Basic Concepts – Shared Memory Segments

1. Shared memory segments are segments of memory that is shared between processes.
2. The steps:
 - 2.1. Create and connect to the shared memory segment, and get a pointer
 - 2.2. Attach to segment
 - 2.3. Read and Write to this pointer and all changes are visible to anyone else connected to the segment.
 - 2.4. Creating the segment and connecting.
3. Creating and connecting to the shared segment

```
int shmget(key_t key, size_t size, int shmflg);
```

where `shmget()` returns an identifier for the shared memory segment,

`key` is hardcoded or created using `ftok()`

`size` size in bytes of shared memory segment.

`shmflg` permissions of segment

- OR'd with `IPC_CREAT` to create segment.
- can be 0.

- `IPC_CREAT` can always be specified.
It will connect, if the segment already exists.

Here's an example to create a 4K segment with 644 permissions (`rw-r--r--`):

```
key_t key = 0xFEEED4BBC;

int shmid;

shmid = shmget( key, 1024, 0644 | IPC_CREAT );
```

4. Attaching to the shared segment

`shmat()` gets a pointer to that data from the `shmid` handle.

Before using a shared memory segment, it needs to be attached using the `shmat()` call:

```
void *shmat(int shmid, void *shmaddr, int shmflg);
```

where

<code>shmget()</code>	returns an identifier for the shared memory segment,
<code>key</code>	is hardcoded or created using <code>ftok()</code>
<code>size</code>	size in bytes of shared memory segment.
<code>shmflg</code>	permissions of segment
	- OR'd with <code>IPC_CREAT</code> to create segment.
	- can be 0.

5. Detaching and deleting shared segments

```
int shmdt(void *shmaddr);
```

`shmaddr`, is the address you got from `shmat()`.

The function returns -1 on error, 0 on success.

6. Sample Code (Sourced From <http://www.ecst.csuchico.edu/~beej/guide/ipc/shmem.html>)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define SHM_SIZE 1024 /* make it a 1K shared memory segment */

int main(int argc, char *argv[])
{
    key_t key;
    int shmid;
    char *data;
```

```
int mode;

if (argc > 2) {
    fprintf(stderr, "usage: shmdemo [data_to_write]\n");
    exit(1);
}

/* make the key: */
if ((key = ftok("shmdemo.c", 'R')) == -1) {
    perror("ftok");
    exit(1);
}

/* connect to (and possibly create) the segment: */
if ((shmid = shmget(key, SHM_SIZE, 0644 | IPC_CREAT)) == -1) {
    perror("shmget");
    exit(1);
}

/* attach to the segment to get a pointer to it: */
data = shmat(shmid, (void *)0, 0);
if (data == (char *)(-1)) {
    perror("shmat");
    exit(1);
}

/* read or modify the segment, based on the command line: */
if (argc == 2) {
    printf("writing to segment: \"%s\"\n", argv[1]);
    strncpy(data, argv[1], SHM_SIZE);
} else
    printf("segment contains: \"%s\"\n", data);

/* detach from the segment: */
if (shmdt(data) == -1) {
    perror("shmdt");
    exit(1);
}

return 0;
}
```

semaphores

semaphore

- Concurrency control mechanism

Slide #8-7

Notes**Basic Concepts – Semaphores**

1. Semaphore is a lock.
It is used for concurrency control.
 - 1.1. Create semaphore set using `semget()`
 - 1.2. and connect to the semaphore set, and get a pointer
 - 1.3. Lock, and access
or, Wait, till semaphore lock is released
 - 1.4. Unlock
2. The steps:
<http://www.ecst.csuchico.edu/~beej/guide/ipc/semaphores.html>

semget()

- 2.1. Creates a semaphore set
- 2.2. Returns an *id* to a semaphore set.
A semaphore set can contain one or many semaphores.

```
#include <sys/sem.h>
```

```
int semget(key_t key, int nsems, int semflg);
```

key is a unique identifier used by different processes to identify this semaphore set.

nsems, is the number of semaphores in this semaphore set.

semflg is the permissions on the new semaphore set

```
#include <sys/ipc.h>
#include <sys/sem.h>

key_t key;
int semid;

key = ftok("somefile", 'E');
semid = semget(key, 10, 0666 | IPC_CREAT);
```

semop()

2.3. set, get, or test-n-set a semaphore

2.4. functionality depends on struct sembuf

```
unsigned short sem_num; /* semaphore number */
short          sem_op;  /* semaphore operation */
short          sem_flg; /* operation flags */
```

sem_op	Description
Positive	used to mark a resource as allocated. value of sem_op is added to semaphore's value. This operation can always proceed. It never forces a process to wait. The calling process must have alter permission on the semaphore set.
Negative	If the absolute value of sem_op is greater than the value of the semaphore, the calling process will block until the value of the semaphore reaches that of the absolute value of sem_op. Finally, the absolute value of sem_op will be subtracted from the semaphore's value. This is how a process releases a resource guarded by the semaphore.
Zero	This process will wait until the semaphore in question reaches 0.

<http://www.ecst.csuchico.edu/~beej/guide/ipc/semaphores.html>

`seminit.c` creates the semaphore. Try using `ipcs` from the command line to verify that it exists. Then run `semdemo.c` in a couple of windows and see how they interact. Finally, use `semrm.c` to remove the semaphore. You could also try removing the semaphore while running `semdemo.c` just to see what kinds of errors are generated.

Here's `seminit.c` (run this first!):

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
```

```

#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int main(void)
{
    key_t key;
    int semid;
    union semun arg;

    if ((key = ftok("semdemo.c", 'J')) == -1) {
        perror("ftok");
        exit(1);
    }

    /* create a semaphore set with 1 semaphore: */
    if ((semid = semget(key, 1, 0666 | IPC_CREAT)) == -1) {
        perror("semget");
        exit(1);
    }

    /* initialize semaphore #0 to 1: */
    arg.val = 1;
    if (semctl(semid, 0, SETVAL, arg) == -1) {
        perror("semctl");
        exit(1);
    }

    return 0;
}

```

Here's semdemo.c:

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int main(void)
{
    key_t key;
    int semid;
    struct sembuf sb = {0, -1, 0}; /* set to allocate resource */

    if ((key = ftok("semdemo.c", 'J')) == -1) {
        perror("ftok");
        exit(1);
    }

    /* grab the semaphore set created by seminit.c: */
    if ((semid = semget(key, 1, 0)) == -1) {
        perror("semget");
        exit(1);
    }

    printf("Press return to lock: ");
    getchar();
}

```



```
printf("Trying to lock...\n");
if (semop(semid, &sb, 1) == -1) {
    perror("semop");
    exit(1);
}
printf("Locked.\n");
printf("Press return to unlock: ");
getchar();
sb.sem_op = 1; /* free resource */
if (semop(semid, &sb, 1) == -1) {
    perror("semop");
    exit(1);
}
printf("Unlocked\n");
return 0;
}
```

Here's [semrm.c](#):

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
int main(void)
{
    key_t key;
    int semid;
    union semun arg;

    if ((key = ftok("semdemo.c", 'J')) == -1) {
        perror("ftok");
        exit(1);
    }

    /* grab the semaphore set created by seminit.c: */
    if ((semid = semget(key, 1, 0)) == -1) {
        perror("semget");
        exit(1);
    }

    /* remove it: */
    if (semctl(semid, 0, IPC_RMID, arg) == -1) {
        perror("semctl");
        exit(1);
    }

    return 0;
}
```

Message Queues

Message queues

Slide #8-8

Notes

Basic Concepts – Message Queues

1. Message Queues are a queue for allowing multiple processes to communicate with each other
2. The steps:

- 2.1. Create or connect to the message queue, and get a pointer

```
int msgget(key_t key, int msgflg);
```

`msgget()` returns the message queue ID on success, or -1 on failure

- 2.2. Send to Queue

Each message is made up of two parts, which are defined in the template structure

`struct msgbuf`, as defined in `sys/msg.h`:

```
struct msgbuf {  
    long mtype;  
    char mtext[1];  
};
```

The field `mtype` is used later when retrieving messages from the queue, and can be set to any positive number. `mtext` is the data this will be added to the queue.

```
int msgsnd(int msqid, const void *msgp, size_t msgsz, int msgflg);
```

`msqid` is the message queue identifier returned by `msgget()`. The pointer `msgp` is a pointer to

the data you want to put on the queue. *msgsz* is the size in bytes of the data to add to the queue. Finally, *msgflg* allows you to set some optional flag parameters,

2.3. Receive from Queue

```
int msgrcv(int msqid, void *msgp, size_t msgsz, long msgtyp, int msgflg);
```

the behavior of `msgrcv()` can be modified by choosing a *msgtyp* that is positive, negative, or zero:

<i>msgtyp</i>	Effect on <code>msgrcv()</code>
Zero	Retrieve the next message on the queue, regardless of its <code>mtype</code> .
Positive	Get the next message with an <code>mtype</code> equal to the specified <i>msgtyp</i> .
Negative	Retrieve the first message on the queue whose <code>mtype</code> field is less than or equal to the absolute value of the <i>msgtyp</i> argument.

2.4. Destroy Queue

```
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
```

`cmd` => `IPC_RMID`, which is used to remove the message queue

```
msgctl(msqid, IPC_RMID, NULL);
```

<http://www.ecst.csuchico.edu/~beej/guide/ipc/mq.html>

```
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

struct my_msgbuf {
    long mtype;
    char mtext[200];
};

int main(void)
{
    struct my_msgbuf buf;
    int msqid;
    key_t key;

    if ((key = ftok("kirk.c", 'B')) == -1) {
        perror("ftok");
        exit(1);
    }

    if ((msqid = msgget(key, 0644 | IPC_CREAT)) == -1) {
        perror("msgget");
        exit(1);
    }

    printf("Enter lines of text, ^D to quit:\n");

    buf.mtype = 1; /* we don't really care in this case */
```

```

    while(gets(buf.mtext), !feof(stdin)) {
        if (msgsnd(msqid, (struct msgbuf *)&buf, sizeof(buf), 0) == -1)
            perror("msgsnd");
    }

    if (msgctl(msqid, IPC_RMID, NULL) == -1) {
        perror("msgctl");
        exit(1);
    }

    return 0;
}

```

and

```

#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

struct my_msgbuf {
    long mtype;
    char mtext[200];
};

int main(void)
{
    struct my_msgbuf buf;
    int msqid;
    key_t key;

    if ((key = ftok("kirk.c", 'B')) == -1) { /* same key as kirk.c */
        perror("ftok");
        exit(1);
    }

    if ((msqid = msgget(key, 0644)) == -1) { /* connect to the queue */
        perror("msgget");
        exit(1);
    }

    printf("spock: ready to receive messages, captain.\n");

    for(;;) { /* Spock never quits! */
        if (msgrcv(msqid, (struct msgbuf *)&buf, sizeof(buf), 0, 0) ==
-1) {
            perror("msgrcv");
            exit(1);
        }
        printf("spock: \"%s\"\n", buf.mtext);
    }

    return 0;
}

```

Chapter Summary

For this chapter, the following are the objectives:

- getopts,
- shared memory segments
- semaphores,
- message queues

Slide #8-9

Chapter 8

IPC - II

Assignment Questions

Questions:

8.1 Modify “mysh” (from chapter 7) to include parsing of command line arguments, such that

- 8.1.1 Using the ‘-pipe’ argument will use the PIPE IPC mechanism
- 8.1.2 Using the ‘-sysv’ argument will use the IPC mechanism of message queues to pass commands; the shared memory segment is used for any communication back to the parent. And, semaphores to regulate the access to the shared memory segments.

Useful links

<http://www.ecst.csuchico.edu/~beej/guide/ipc/shmem.html>

<http://www.ecst.csuchico.edu/~beej/guide/ipc/mq.html#ftok>

<http://www.ecst.csuchico.edu/~beej/guide/ipc/semaphores.html>